

Spam Detection with Machine Learning

Machine Learning

Arthur Samuel, an early American leader in the field of computer gaming and artificial intelligence, coined the term "Machine Learning" in 1959 while at IBM. He defined machine learning as "the field of study that gives computers the ability to learn without being explicitly programmed".

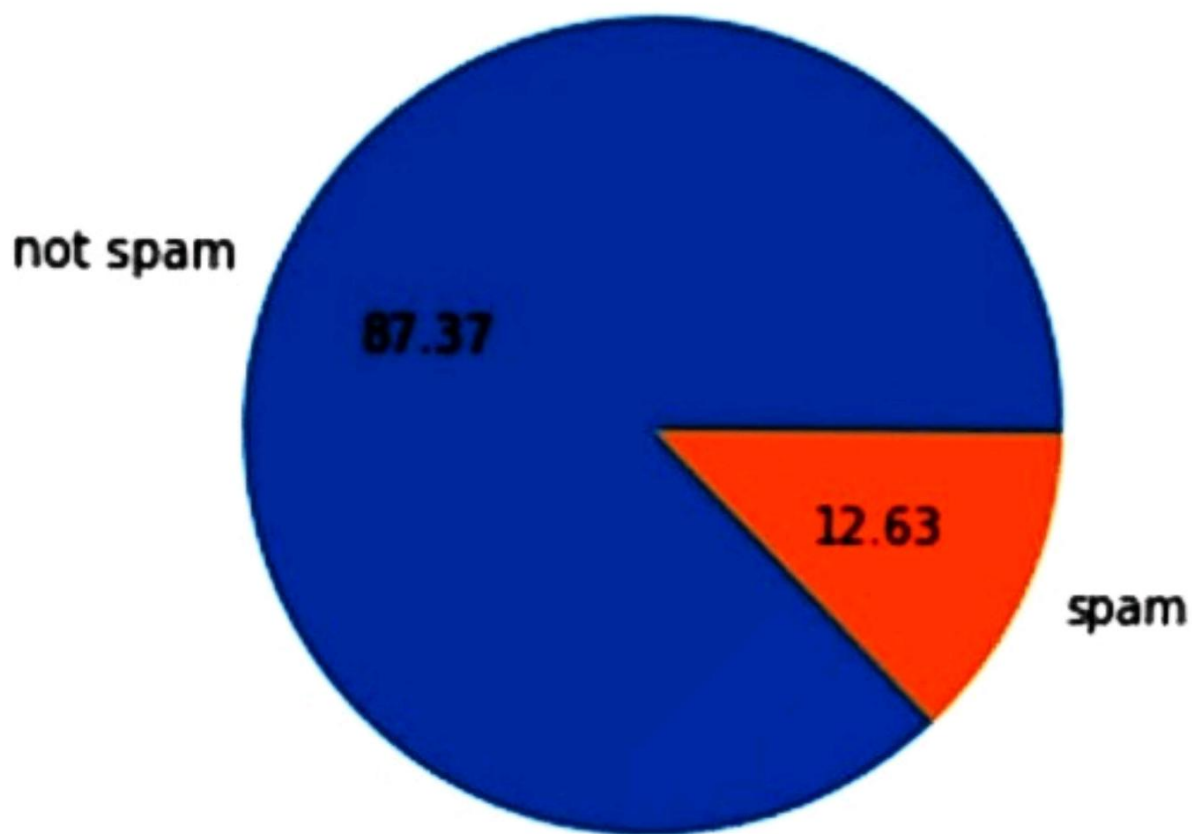
- Machine learning is programming computers to optimize a performance criterion using example data or past experience.
- The field of study known as machine learning is concerned with the question of how to construct computer programs that automatically improve with experience.

Naïve Bayes

- Use Bayes Theorem: $P(H|e) = \frac{P(H|e)P(e)}{P(H)}$
- Hypothesis (H): spam or not spam
- Event (e): word occurs
- For example, the probability an email is spam when the word "free" is in the email

$$P(spam | "free") = \frac{P("free" | spam)P(spam)}{P("free")}$$

- "Naïve": assume the feature values are independent of each other



```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import
train_test_split
from sklearn.feature_extraction.text
import TfidfVectorizer
from sklearn.naive_bayes import
MultinomialNB
from sklearn.metrics import
accuracy_score, confusion_matrix,
classification_report

# Load your dataset (assuming you
have a CSV file with 'text' and 'label'
columns)
data = pd.read_csv('spam_data.csv')

# Split the data into training and testing
sets
X = data['text']
```

```
y = data['label']
X_train, X_test, y_train, y_test =
train_test_split(X, y, test_size=0.2,
random_state=42)

# Text vectorization using TF-IDF
tfidf_vectorizer = TfidfVectorizer()
X_train_tfidf =
tfidf_vectorizer.fit_transform(X_train)
X_test_tfidf =
tfidf_vectorizer.transform(X_test)

# Train a Naive Bayes classifier
spam_classifier = MultinomialNB()
spam_classifier.fit(X_train_tfidf, y_train)

# Make predictions on the test data
y_pred =
spam_classifier.predict(X_test_tfidf)
```

```
# Evaluate the classifier
accuracy = accuracy_score(y_test,
y_pred)
confusion = confusion_matrix(y_test,
y_pred)
report = classification_report(y_test,
y_pred)

print(f'Accuracy: {accuracy}')
print('Confusion Matrix:')
print(confusion)
print('Classification Report:')
print(report)
...
```

Accuracy: 0.965

Confusion Matrix:

```
[[958   7]
 [ 26 134]]
```

Classification Report:

		precision	recall
f1-score	support		
	ham	0.97	0.99
0.98	965		
	spam	0.95	0.84
0.89	160		
	accuracy		
0.97	1125		
	macro avg	0.96	0.91
0.93	1125		
	weighted avg	0.97	0.97
0.97	1125		

Conclusion

- ❑ Legitimacy Score
 - No content needed
- ❑ Can Be Combined with Content-Based Filters
- ❑ More Sophisticated Classifiers
 - SVM, boosting, etc
- ❑ Classifiers Using Combined Feature